

Main.py

```
1 import pandas as pd
2 import numpy as np
3
4 # Define paths for each dataset in your Google Drive
5 df1_path = '/Users/ankushbhatt/Downloads/VFS-CS/DataSet/groundtruth.csv' # Update
   with your actual path
6 df2_path = '/Users/ankushbhatt/Downloads/VFS-CS/DataSet/weather_features.csv' #
   Update with your actual path
7 df3_path = '/Users/ankushbhatt/Downloads/VFS-CS/DataSet/road_features.csv' # Update
   with your actual path
8
9 # define the datatypes for each csv
10 dtypes_df1 = {'road_segment_id': 'int32', 'timestamp': 'str', 'max_capacity':
   'int32', 'occupied': 'int32', 'available': 'int32'}
11 dtypes_df2 = {'road_segment_id': 'int32', 'timestamp': 'str', 'tempC': 'float32',
   'windspeedKmph': 'float32', 'precipMM': 'float32'}
12 dtypes_df3 = {'road_segment_id': 'int32', 'commercial': 'Int64', 'residential':
   'Int64', 'transportation': 'Int64', 'schools': 'Int64', 'eventsites': 'Int64',
13               'restaurant': 'Int64', 'shopping': 'Int64', 'office': 'Int64',
   'supermarket': 'Int64', 'num_off_street_parking': 'Int64', 'off_street_capa':
   'Int64'}
14
15 # load the datasets using the defined paths
16 df1 = pd.read_csv(df1_path, dtype=dtypes_df1)
17 df2 = pd.read_csv(df2_path, dtype=dtypes_df2)
18 df3 = pd.read_csv(df3_path, dtype=dtypes_df3)
19
20
21 # merge the dataset
22 # first merge df1 and df2 on road_segment_id and timestamp
23 merged_df = pd.merge(df1, df2, on=['road_segment_id', 'timestamp'], how='inner')
24 # second merge df3 with merged_df on road_segment_id
25 final_df = pd.merge(merged_df, df3, on='road_segment_id', how='inner')
26
27 # save the merged dataset
28 final_df.to_csv('New_joint_dataset.csv', index=False)
29
30 # Print the final merged dataset info
31 print("\nMerged Dataset:")
32 print(final_df.info()) # showing information about final merged dataset
33 print(final_df.head()) # showing first few rows of finam merged dataset
34
35 # Data Pre processing
36
37 # Load New Joint dataset
38
39 import pandas as pd
40 import numpy as np
41 df = pd.read_csv('New_joint_dataset.csv')
42
43 # get a summary of the data
44 print("First few rows of dataframe")
```

```
45 print(df.head())
46 print("\ndataframe info:")
47 print(df.info())
48 print("\nsummary of dataframe:")
49 print(df.describe(include='all'))
50
51 # Check for missing values
52 missing_values = df.isnull().sum()
53 print("\nMissing Values:")
54 print(missing_values)
55
56 # Fill the missing values
57 # Now fill numeric values with median and categorical with mode
58 for column in final_df.columns:
59     if df[column].dtype in ['float64', 'int64']:
60         df[column].fillna(df[column].median(), inplace=True)
61     elif df[column].dtype == 'object':
62         df[column].fillna(df[column].mode()[0], inplace=True)
63
64 # Identify and handle the outliers
65 # Now using IQR methods for detecting outliers in the data
66 def remove_outliers_iqr(data):
67     for column in data.select_dtypes(include=['float64', 'int64']).columns:
68         Q1 = data[column].quantile(0.25)
69         Q3 = data[column].quantile(0.75)
70         IQR = Q3 - Q1
71         lower_bound = Q1 - 1.5 * IQR
72         upper_bound = Q3 + 1.5 * IQR
73         data = data[(data[column] >= lower_bound) & (data[column] <= upper_bound)]
74     return data
75 df = remove_outliers_iqr(df)
76
77 # Standardize the format of data
78 df['timestamp'] = pd.to_datetime(df['timestamp'], errors='coerce') # 'coerce' will
turn invalid parsing into NaT (Not a Time)
79
80 # Extract useful date-time features
81 df['year'] = df['timestamp'].dt.year
82 df['month'] = df['timestamp'].dt.month
83 df['day'] = df['timestamp'].dt.day
84 df['hour'] = df['timestamp'].dt.hour
85 df['minute'] = df['timestamp'].dt.minute
86 df['dayofweek'] = df['timestamp'].dt.dayofweek # Monday=0, Sunday=6
87
88 # Drop the original 'timestamp' column as it is no longer needed for prediction
89 df.drop(columns=['timestamp'], inplace=True)
90
91 # Now you can proceed with other steps (e.g., missing values handling, outliers
removal, etc.)
92
93 # Now removing duplicates
94 df.drop_duplicates(inplace=True)
95
96 # dropping the unnecessary columns
```

```
97 df.drop(columns=['Unnamed: 0_x', 'Unnamed: 0_y', 'Unnamed: 0', 'max_capacity',
98 'occupied'], inplace=True, errors='ignore')
99
100 # convert data types
101 # convert numeric columns to appropriate types
102 from sklearn.preprocessing import StandardScaler, PowerTransformer
103 numeric_columns = ['max_capacity', 'occupied', 'available', 'tempC',
104 'windspeedKmph', 'precipMM', 'commercial', 'residential', 'transportation',
105 'schools', 'eventsites', 'restaurant', 'shopping',
106 'office', 'supermarket', 'num_off_street_parking',
107 'off_street_capa']
108
109 for column in numeric_columns:
110     if column in df.columns:
111         df[column] = pd.to_numeric(df[column], errors='coerce')
112
113 # convert 'road_segment_id' to int32 if necessary
114 if 'road_segment_id' in df.columns:
115     df['road_segment_id'] = df['road_segment_id'].astype('int32')
116
117 ***Skewness Correction with Power Transformation**
118 # Apply PowerTransformer to correct skewness in numeric columns
119 numeric_features = df.select_dtypes(include=['float64', 'int64']).columns
120 power_transformer = PowerTransformer(method='yeo-johnson', standardize=True)
121 df[numeric_features] = power_transformer.fit_transform(df[numeric_features])
122
123 # **Scaling Features**
124 # Use StandardScaler to standardize numeric features
125 scaler = StandardScaler()
126 df[numeric_features] = scaler.fit_transform(df[numeric_features])
127
128 # check datatypes after conversion
129 print("\nData Types after Conversion:")
130 print(df.dtypes)
131
132 # check the info of the final dataframe
133 print("\nFinal Dataframe Info:")
134 print(df.info())
135
136 # save the clean data
137 clean_data_path = 'clean_data.csv'
138 df.to_csv(clean_data_path, index=False)
139 print("\nData preprocessing completed. clean data saved to", clean_data_path)
140
141 # Trained the Model:
142 import pandas as pd
143 import numpy as np
144 from sklearn.model_selection import train_test_split, RandomizedSearchCV,
145 cross_val_score
146 from sklearn.ensemble import RandomForestRegressor
147 from sklearn.metrics import mean_absolute_error, mean_squared_error
148 import matplotlib.pyplot as plt
149
150 # Load the cleaned dataset
```

```
146 cd = pd.read_csv('clean_data.csv')
147
148 # Features and Target
149 x = cd.drop(columns=['available']) # Features
150 y = cd['available'] # Target
151
152 # Split the data into train and test sets (80% train, 20% test)
153 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
154 random_state=42)
155
156 # Define the model with slightly fewer trees and shallower depth
157 model = RandomForestRegressor(
158     n_estimators=30, # Reduced n_estimators for faster learning
159     max_depth=5, # Limiting tree depth to avoid overfitting
160     min_samples_split=5, # Minimum samples required to split a node
161     min_samples_leaf=4, # Minimum samples required to be in a leaf node
162     random_state=42,
163     n_jobs=-1 # Use all available CPU cores for faster training
164 )
165
166 # Fit the model on training data
167 model.fit(x_train, y_train)
168
169 # Predict on test data
170 y_pred = model.predict(x_test)
171
172 # Evaluate the model
173 mae = mean_absolute_error(y_test, y_pred)
174 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
175 r2_score = model.score(x_test, y_test)
176
177 # Print evaluation metrics
178 print("Model Evaluation Metrics:")
179 print("Mean Absolute Error (MAE):", mae)
180 print("Root Mean Squared Error (RMSE):", rmse)
181 print("R² (Coefficient of Determination) accuracy:", r2_score)
182
183 # Cross-validation with fewer folds (3-fold to save time)
184 cv_scores = cross_val_score(model, x, y, cv=5, scoring='neg_mean_squared_error',
185 n_jobs=-1)
186 print("\nCross-validation scores:", cv_scores)
187 print("Average Cross-validation score:", cv_scores.mean())
188
189 # Hyperparameter tuning with a focused parameter grid and reduced combinations
190 param_distributions = {
191     'n_estimators': [10, 30, 50], # Reduced range for faster tuning
192     'max_depth': [5, 7], # Limiting to fewer options
193     'min_samples_split': [2, 5], # Narrowing range
194     'min_samples_leaf': [1, 2, 4], # Fewer leaf node options
195 }
196
197 # RandomizedSearchCV with fewer iterations
198 randomized_search = RandomizedSearchCV(
199     RandomForestRegressor(random_state=42, n_jobs=-1),
```

```
198     param_distributions,
199     n_iter=5,                # Fewer combinations to try
200     cv=5,                   # Reduced folds
201     scoring='neg_mean_squared_error',
202     n_jobs=-1
203 )
204
205 # Fit RandomizedSearchCV to training data
206 randomized_search.fit(x_train, y_train)
207
208 # Print best hyperparameters from RandomizedSearchCV
209 print("Best Hyperparameters from RandomizedSearchCV:")
210 print(randomized_search.best_params_)
211
212 # Retrain the model with the best parameters
213 best_model = randomized_search.best_estimator_
214 best_model.fit(x_train, y_train)
215
216 # Predict and evaluate the retrained model
217 y_pred_best = best_model.predict(x_test)
218
219 # Evaluate the model
220 mae_best = mean_absolute_error(y_test, y_pred_best)
221 rmse_best = np.sqrt(mean_squared_error(y_test, y_pred_best))
222 r2_score_best = best_model.score(x_test, y_test)
223
224 # Print evaluation metrics of the best model
225 print("\nBest Model Evaluation Metrics:")
226 print("Mean Absolute Error (MAE):", mae_best)
227 print("Root Mean Squared Error (RMSE):", rmse_best)
228 print("R² (Coefficient of Determination) accuracy:", r2_score_best)
229
230 # Final model accuracy
231 accuracy_percentage_best = r2_score_best * 100
232 print(f"Final Accuracy: {accuracy_percentage_best:.2f}%")
233
234 # Cross-validation to evaluate model performance on multiple folds (using 3-fold CV)
235 cv_scores_best = cross_val_score(best_model, x, y, cv=5,
236                                   scoring='neg_mean_squared_error', n_jobs=-1)
237 print("\nCross-validation scores of best model:", cv_scores_best)
238 print("Average Cross-validation score of best model:", cv_scores_best.mean())
```